


# 2022 FALL OS Project 1 Help Document

---



**JINYONG HA**

**Distributed Computing Systems Laboratory  
Department of Computer Science and Engineering  
Seoul National University, Korea**

**2022-09-23**

# NOTE

---

- 과제 **#0** Set up development environment

기한: 09/17 midnight

방법: etl 제출

- 과제 **#1**

기한: 10/08 midnight

방법: 팀별 repository에 proj1 branch 생성 / 제출물은 대표 한 명이 etl에

- 과제 **#2**

기한: 11/05 midnight

방법: 팀별 repository에 proj2 branch 생성 / 제출물은 대표 한 명이 etl에

# Github repository

[main] <https://github.com/jyha200/osfall2022>

[per-team] [https://github.com/\[github\\_id\]/osfall2022-team\[#\]](https://github.com/[github_id]/osfall2022-team[#])

reminder:

아래 링크에 **Github ID**를 입력해주세요

<https://docs.google.com/spreadsheets/d/1VBGmW-qkLfKm7Felp5JhEDqcJ0CYFrX3/edit?usp=sharing&oid=106917183581249491181&rtpof=true&sd=true>

# What you did in project 0

1. Build your kernel
2. Make image files
3. emulate Tizen with QEMU

# What is QEMU

QEMU is a generic and open source machine emulator and virtualizer.

## 1. Install QEMU

```
sudo apt-get install qemu (or sudo apt-get install  
qemu-system-aarch64)
```

## 2. Update /arch/arm64/configs/tizen\_bcmrpi3\_defconfig file

## 3. Run Tizen on RPI3 emulation with QEMU

5

# Move files into Tizen when using QEMU

1. Mount `rootfs.img` on `${mnt_dir}`
2. Move files under `${mnt_dir}/root/` (You may need sudo)
3. Unmount `${mnt_dir}`

---

- Project 1

# General Overview of Project 1

- Write a system call **as a kernel module**

- `int ptree ( struct prinfo *buf, int *nr)`
- System call number **398**
- You can name your function `sys_ptree`; doesn't matter as long as it works
- **Tizen repository for project 1 updated**
  - Pull & checkout tizen git repository
    - [https://github.com/jyha200/tizen-5.0-rpi3/tree/project1\\_base](https://github.com/jyha200/tizen-5.0-rpi3/tree/project1_base)
  - Contains
    - Simple kernel module example (`kernel/hello_mod.c`)
    - Some changes to ease project

- Test your system call

- Print the entire process tree in pre-order



# Example Program Output

- swapper/0 (pid 0)
  - The first ever process created
  - Used to represent the state of 'not working'
- systemd (pid 1)
  - Manages all the processes
- kthreadd (pid 2)
  - Kernel thread daemon
  - kthread\_create

```
sh-3.2# ./proj1
swapper/0,0,0,0,1,0,0
    systemd,1,1,0,167,2,0
        systemd-journal,167,1,1,0,185,0
        systemd-udevd,185,1,1,0,241,0
        dbus-daemon,241,1,1,0,297,81
        amd,297,1,1,0,298,301
        dlog_logger,298,1,1,0,307,1901
        buxton2d,307,1,1,0,313,375
        key-manager,313,1,1,0,325,444
```

```
    kthreadd,2,1,0,3,0,0
        kworker/0:0,3,1026,2,0,4,0
        kworker/0:0H,4,1026,2,0,5,0
        kworker/u8:0,5,1026,2,0,6,0
        mm_percpu_wq,6,1026,2,0,7,0
        ksoftirqd/0,7,1,2,0,8,0
        rcu_preempt,8,1026,2,0,9,0
        rcu_sched,9,1026,2,0,10,0
```

# Return Value

- Success

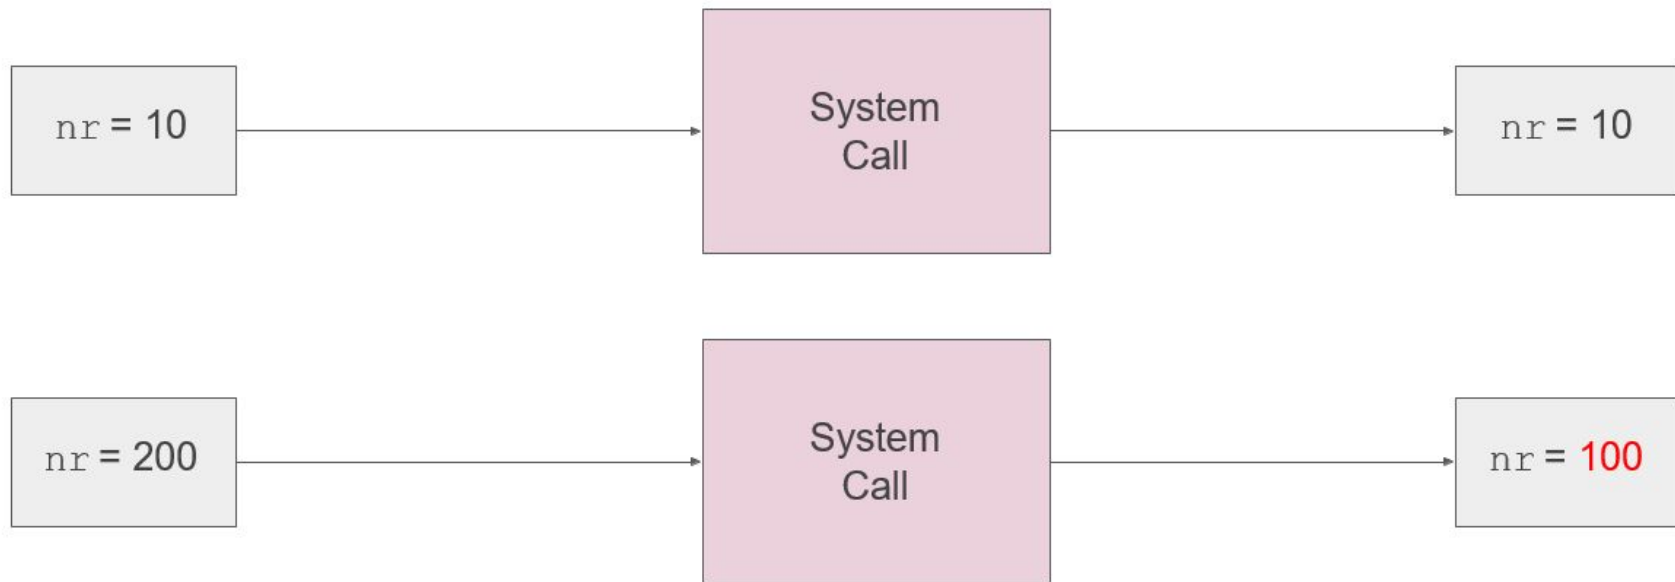
- Your system call should return the total number of entries (this may be bigger than the actual number of entries copied).
- `nr` can be changed

- Error

- Error handling: `-EINVAL` or `-EFAULT`
- You may handle other errors but we will not grade them
- Defined in `include/uapi/asm-generic/errno-base.h`

# nr EXAMPLE

- Assuming the number of total processes is 100:



# Error Handling

## -EINVAL

- If `buf` and/or `nr` are `NULL`, or if `nr` is less than 1

## -EFAULT

- If `buf` and/or `nr` are outside the accessible address space

## ●How to print error messages?

- `int result = syscall(398, ...);`
- `printf("%d", result);`
- `??????`

## ●You cannot get `-EINVAL` or `-EFAULT` as the return value

- Use `errno` and `perror()`

# Check Before Submission!!!!

- **Implement your own code**
  - You might refer to others code but implementation should be made by yourself with your own style.
  - **Copy 발견 시 불이익 있을 수 있음**
- **All your changes should be included in 'kernel/ptree\_mod.c'**
- **Unsafe access** to user space memory
- Return value
  - Incorrect return value
  - Not modifying `*nr` when needed
- Whether you follow the project specifications (final check!) ...
- ... and whether you have delineated all unspecified/different implementation
  - details in README
- **White-box test for this project!**

# Reminder

## **concise README file**

- Describe how to build your kernel
- Describe the high-level design and implementation
- Investigation of the process tree
- Any lessons learned

## **concise 4-minute presentation slides**(including video demo submitted on etl)

- Limit: 10 slides including title slide
- We will not consult slides after 10

Presentation should include

- A. High level design and implementation
- B. video clip that shows that your system works
- C. investigation of the process trees
- D. lesson learned

# About Submission (IMPORTANT!)

- Make sure your branch name is *proj1*
- **Cheating might be penalized**
- **Don't be late! marks will be deducted**

We allow 3 days late submission! TA will not grade any submissions after that!

- Slides and Demo

Submit to etl before deadline.

we allow 3 day late submission, after that submission on etl will not be possible

- zip dir title: [OS-ProjX] TeamX submission

inside the directory there should be 4 files:

- The slides to your presentation: TeamX-slides.ppt(.pdf)
- Your Demo Video: TeamX-demo.mp4(.avi....)
- Recorded presentation(sound only): TeamX-presentation.mp3(.wav....)
- ptree\_mod.c (for verification)
  - **should be buildable when just place submitted ptree\_mod.c to 'kernel' of project 1 base branch ([https://github.com/jyha200/tizen-5.0-rpi3/tree/project1\\_base](https://github.com/jyha200/tizen-5.0-rpi3/tree/project1_base))**

---

# Kernel Programming Guideline



# Important Directories

`arch`

Architecture dependent (i.e. x86, arm, mips, ...) parts of Linux

`kernel`

Common kernel code

`net`

Common network related code

`drivers`

Common driver code for Linux

`fs`

Common file system code for Linux

`include`

Common header files

# Things to Keep in Mind...

- No memory protection
  - Corruption in kernel memory space can make the whole machine crash!
- No floating point or MMX operation
  - Dealing with real numbers can be challenging and painful!
- You unfortunately have to do it for some projects :(
- Rigid stack limit
  - Use extra caution when allocating local arrays or having recursive calls
  - `kmalloc` instead for huge arrays
- Your kernel code will run in a multi-core environment
  - Use proper synchronization mechanisms to avoid race conditions
  - Beware of deadlocks

# Accessing User Memory

- In kernel mode, you should avoid directly accessing user memory space
  - Can result in kernel panic
- `include/asm/uaccess.h` provides macros for this
  - `get_user` / `put_user`: copies simple variables
  - `copy_from_user` / `copy_to_user`: copies a block of data
  - More on <http://www.ibm.com/developerworks/library/l-kernel-memory-access/>
- Mark system call parameters that access user space memory with `__user`
  - e.g. In `include/linux/syscalls.h`:
  - `asmlinkage long sys_time(time_t __user *tloc);`

# The User Space Memory Access API

Function	Description
<code>access_ok</code>	Checks the validity of the user space memory pointer
<code>get_user</code>	Gets a simple variable from user space
<code>put_user</code>	Puts a simple variable to user space
<code>clear_user</code>	Clears, or zeros, a block in user space
<code>copy_to_user</code>	Copies a block of data from the kernel to user space
<code>copy_from_user</code>	Copies a block of data from user space to the kernel
<code>strlen_user</code>	Gets the size of a string buffer in user space
<code>strncpy_from_user</code>	Copies a string from user space into the kernel

# kmalloc and kfree

- Used for allocating / releasing kernel memory instead of `malloc` / `free`
- `kmalloc` is similar to `malloc`, but has an additional flag parameter

```
void *kmalloc(size_t size, int flags)
```

## Frequently used flags

`GFP_KERNEL` Allocate kernel space memory

`GFP_USER` Allocate user space memory

`GFP_ATOMIC` similar to `GFP_KERNEL` , but cannot sleep; used inside interrupts or other non-sleep routines

More on <http://www.makelinux.net/ldd3/chp-8-sect-1.shtml>

- `kfree` is similar to `free`

# task\_struct

- 598 lines!

You don't need to read everything

- children and sibling: doubly linked lists

```
/* Real parent process: */
struct task_struct __rcu          *real_parent;

/* Recipient of SIGCHLD, wait4() reports: */
struct task_struct __rcu          *parent;

/*
 * Children/sibling form the list of natural children:
*/
struct list_head                  children;
struct list_head                  sibling;
struct task_struct                *group_leader;
```

# Doubly Linked List in Linux Kernel

- Linux kernel has a doubly linked list implementation for kernel programming

Extensively used across all Linux kernel code

- Defined in `include/linux/list.h`


Can only be used in kernel space!

- Unlike other commonly used linked lists, kernel list nodes are stored ***inside*** data

```
struct student {  
    char* name;  
    char* student_id; str  
    struct list_head list;  
};
```

# Doubly Linked List in Linux Kernel (cont'd)

```
struct list_head {  
    struct list_head *next, *prev;  
}
```





# Doubly Linked List in Linux Kernel (cont'd)

- Initializing a list node (must be declared beforehand)

- `INIT_LIST_HEAD(&first_student->list)`

- Defining and initializing a list head pointer (declaration + `INIT_LIST_HEAD`)

- `LIST_HEAD(student_list)`

- More about Linux kernel list (highly recommended)

<http://www.makelinux.net/ldd3/chp-11-sect-5.shtml>

# Doubly Linked List in Linux Kernel (cont'd)

- Commonly used macros/functions

`list_add / list_add_tail`: adds a node to a list

`list_del / list_del_init`: deletes a node from a list

`list_for_each_entry`: iterates over a list

`list_for_each_entry_safe`: iterates over a list **when nodes can be deleted**

`list_entry / container_of`: returns the item given a list node

# Kernel Module

---

- What is kernel module?
  - Pluggable binary for extra functionalities of kernel
  - Named as **mod\_name.ko**
- How to use
  - Insert
    - # insmod XX.ko
  - Remove
    - # rmmod XX.ko

# Replace Syscall Table

- Syscall Table
  - Array of syscall function pointers
  - Indexed by syscall ID
  - See `arch/arm64/kernel/sys32.c`
- **Change only 398th entry**
  - Otherwise, it might bother testing
  - May be reflected to your score



# Implement Kernel Module (See kernel/hello\_mod.c)

- `module_init(f1)`
  - Call module initialization function **f1** when module is inserted
- `module_exit(f2)`
  - Call module cleanup function **f2** when module is removed
- `MODULE_LICENSE()`
  - Indicate the license of module
  - Use “GPL v2”

# Useful References

- Linux cross reference (LXR)

<https://elixir.bootlin.com/linux/v4.14.67/source>

- Unreliable Guide To Hacking The Linux Kernel by Rusty Russel

<http://kernelbook.sourceforge.net/kernel-hacking.pdf>

